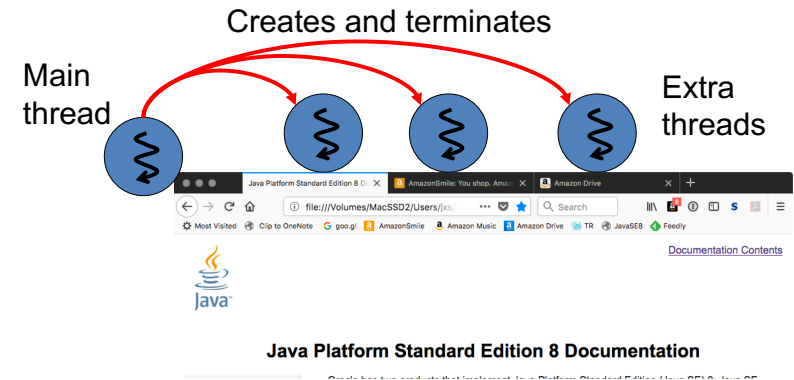


Race Conditions and Thread Synchronization (Locking)

Goals of Concurrency/Multi-threading

- **Responsiveness**



- **Efficiency**

- c.f. MCTest, multi-threaded prime number generation

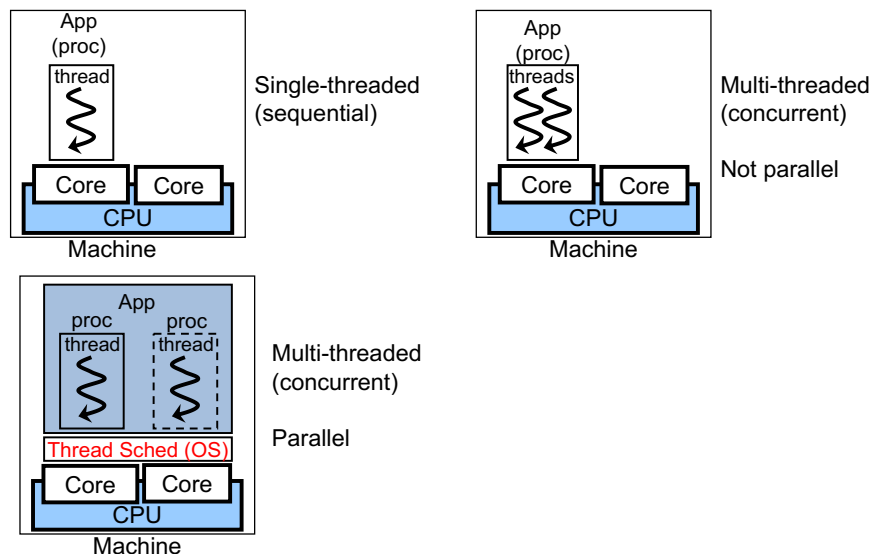
Thread Safety

- Threads are a powerful tool to make your code more responsive and efficient.
- However, multi-threaded code can raise **“thread safety” issues** if it is poorly written.
- 2 major thread safety issues
 - Race conditions (data races)
 - Mess up the consistency of data shared among threads
 - Deadlock
 - Make code execution stuck.
- **Thread-safe code** is free from those 2 issues.

Race Conditions (a.k.a. Data Races)

- Threads run **independently**.
 - No coordination among threads by default.
 - c.f. MCTest, PrimeNumberGenerator
 - join() allows threads to coordinate with each other.
- They can **share** objects/data.
 - Exception: **Local variables** are NOT shared among threads.
- They can **mess up** the consistency of the shared objects/data.
 - A thread can write some data to a variable when another thread is reading data from the variable.
 - A thread can write some data to a variable when another thread is writing different data to the variable.


Single- and Multi-threaded Programs



5


In CS681...

- We always assume a single CPU core that runs multiple threads.
 - The most conservative scenario.
- If your code is thread-safe in the most conservative scenario, it is always thread-safe in less conservative scenarios as well.



An Example Race Condition:

ThreadUnsafeBankAccount.java




- The variable "balance" is shared by 2 threads.
- They access the variable independently.
- ```

public void deposit(double amount){
 System.out.print("Current balance (d): " + balance);
 double newBalance = balance + amount;
 System.out.println("New balance (d): " + newBalance);
 balance = newBalance;
}

public void withdraw(double amount){
 System.out.print("Current balance (w): " + balance);
 double newBalance = balance - amount;
 System.out.println("New balance (w): " + newBalance);
 balance = newBalance;
}

```

7

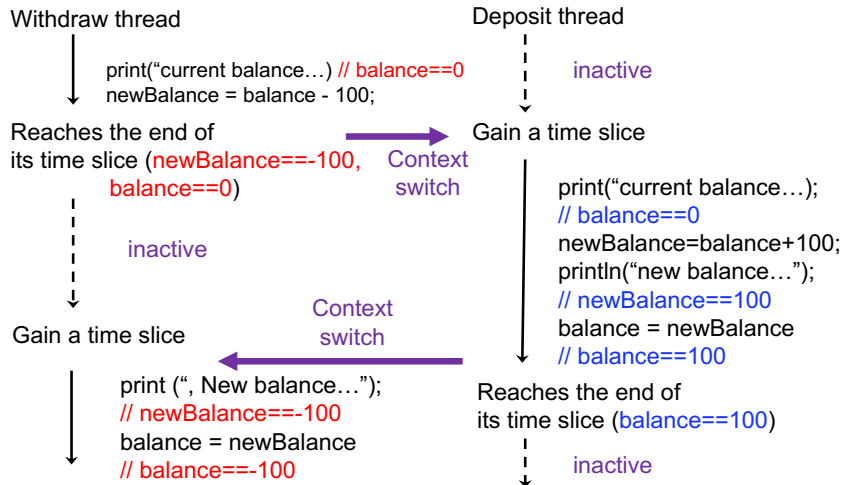


- Desirable output:
  - Current balance (w): 0.0, New balance (w): -100.0
  - Current balance (d): -100.0, New balance (d): 0.0
  - Current balance (w): 0.0, New balance (w): -100.0
  - Current balance (d): -100.0, New balance (d): 0.0
  - Current balance (d): 0.0, New balance (d): 100.0
  - Current balance (w): 100.0, New balance (w): 0.0
  - ...
- In reality:
  - Current balance (w): 0.0, Current balance (d): 0.0, New balance (d): 100.0
  - , New balance (w): -100.0

8

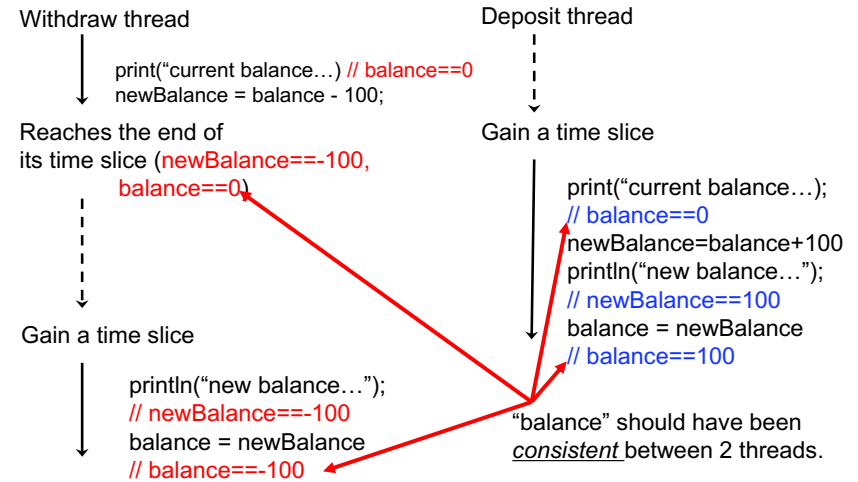
# How Can This Happen?

- Current balance (w): 0.0 Current balance (d): 0.0, New balance (d): 100.0
- , New balance (w): -100.0



9

- Current balance (w): 0.0 Current balance (d): 0.0, New balance (d): 100.0
- , New balance (w): -100.0



10

## The Source of the Problem: Visibility


- ThreadUnsafeBankAccount.java is **NOT thread safe**.
  - Race conditions can occur.
- Race conditions occur due to **visibility** issues.
  - The current (most up-to-date) value of the shared variable (e.g. "balance") is **not visible** for all threads.

## Race Conditions (a.k.a. Data Races)


- All threads
  - Run in their **race** to complete their tasks.
  - Manipulate a shared object/data **independently**.
- The end result depends on which of them happens to win the race.
  - No guarantees on the order of thread execution.
  - No guarantees on how many tasks a thread can perform in a single CPU time slice/quota.
  - No guarantees on the end result on shared data.

## Note: Thread States

- Both **“active”** and **“inactive”** threads are in the **Runnable** state.
  - The **Runnable** state does NOT distinguish if a thread is **“actively running”** on a CPU core or it is **“inactive waiting”** for its next turn.
  - The **Waiting** state does NOT mean that a thread is runnable but inactive.



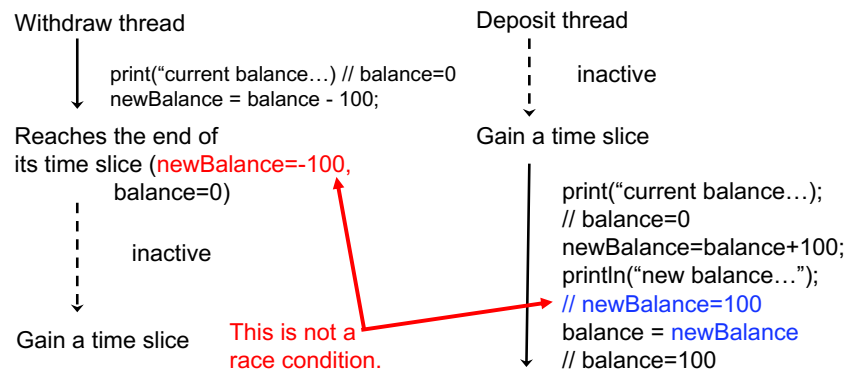
## States of a Thread



14

## Note: Local Variables

- Local variables are NOT shared by threads.
  - It is created and maintained in a **thread-by-thread** manner.
    - The “withdraw” thread has no access to a value of `newBalance` that the “deposit” thread has created.
    - The “deposit” thread has no access to a value of `newBalance` that the “withdraw” thread has created.



- Race conditions never occur due to local variables.
- Focus on non-local (i.e. **shared**) variables in debugging threaded code.
- Two other example local variables

```

- public void deposit(double amount){
 System.out.print("Current balance (d): " + balance);
 double newBalance = balance + amount;
 System.out.println(", New balance (d): " + newBalance);
 balance = newBalance;
}

- public void withdraw(double amount){
 System.out.print("Current balance (w): " + balance);
 double newBalance = balance - amount;
 System.out.println(", New balance (w): " + newBalance);
 balance = newBalance;
}

```

## Another Example:

### ThreadUnsafeBankAccount2

- Local variables (newBalance) are removed from ThreadUnsafeBankAccount

```
- public void deposit(double amount){
 balance = balance + amount;
}
- public void withdraw(double amount){
 balance = balance - amount;
}
```

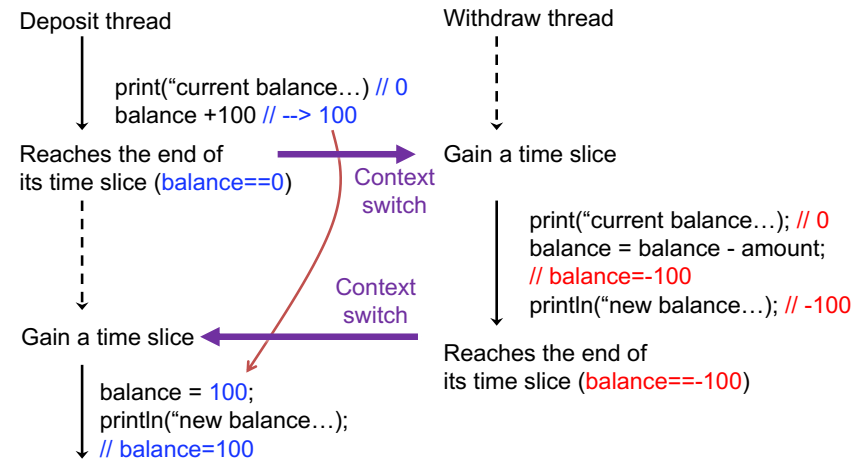
### Output

```
- Current balance (d): 0.0, New balance (d): 100.0
- Current balance (w): 100.0, New balance (w): 0.0
- Current balance (d): 0.0, New balance (d): 100.0
- Current balance (w): 100.0, New balance (w): 0.0
- Current balance (d): 0.0Current balance (w): 0.0, New balance (w):
-100.0
- , New balance (d): 100.0
```

17

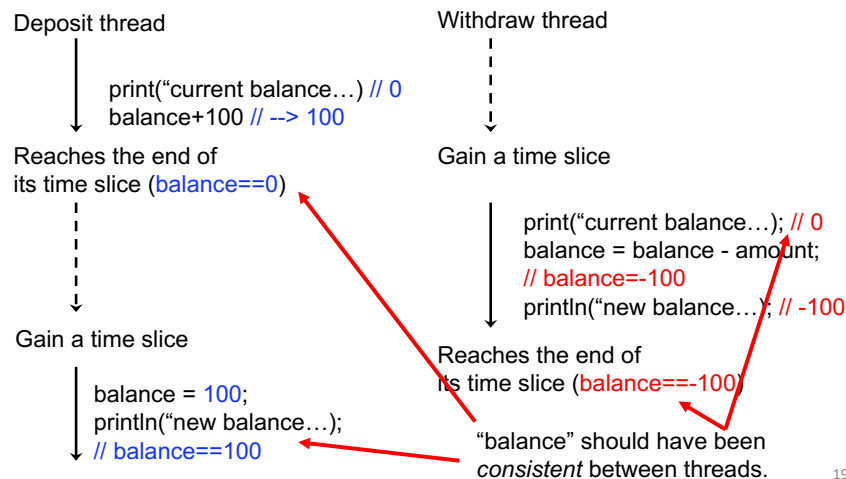
## How Can This Happen?

- Current balance (d): 0.0Current balance (w): 0.0, New balance (w): -100.0
- , New balance (d): 100.0



18

- Current balance (d): 0.0Current balance (w): 0.0, New balance (w): -100.0
- , New balance (d): 100.0



19


- This is not a solution: `balance -= amount;`
  - Just a **syntactic sugar** for `balance = balance - amount;`
- All threads
  - Run in their *race* to complete their tasks.
  - Manipulate a shared object/data independently.
- The end result depends on which of them happens to win the race.
  - No guarantees on the order of thread execution.
  - No guarantees on how many tasks a thread can perform in a single CPU time slice.
  - No guarantees on the end result on shared data.

20


# Solution: Thread Synchronization

- Need to synchronize threads
  - i.e., Need to serialize their concurrent access to a shared variable
- Thread synchronization
  - Enables **serialized** (or **mutually-exclusive**) access to a shared variable
  - Allows **only one thread** to access a shared variable at a time
    - Forces all other threads to wait and take turn to access it.

21



- Thread synchronization
  - Prevents the “**withdraw**” thread from withdrawing money from “balance” when the “**deposit**” thread is depositing money to “balance.”
  - Prevents the “**deposit**” thread from depositing money to “balance” when the “**withdraw**” thread is withdrawing money from “balance.”



## Thread Synchronization with Java

- Java implements **thread synchronization** by
  - Providing **locks**
  - Allowing you to write **atomic code** (a.k.a critical section) with locks.
    - **Atomic code**: A piece of code that is executed by multiple threads in a synchronized (or serialized, or mutually-excluded) manner.
      - When a thread is running atomic code, no other threads can run it.
  - No intermediate results/states produced in atomic code can be revealed/exposed to other threads.

## Locks in Java

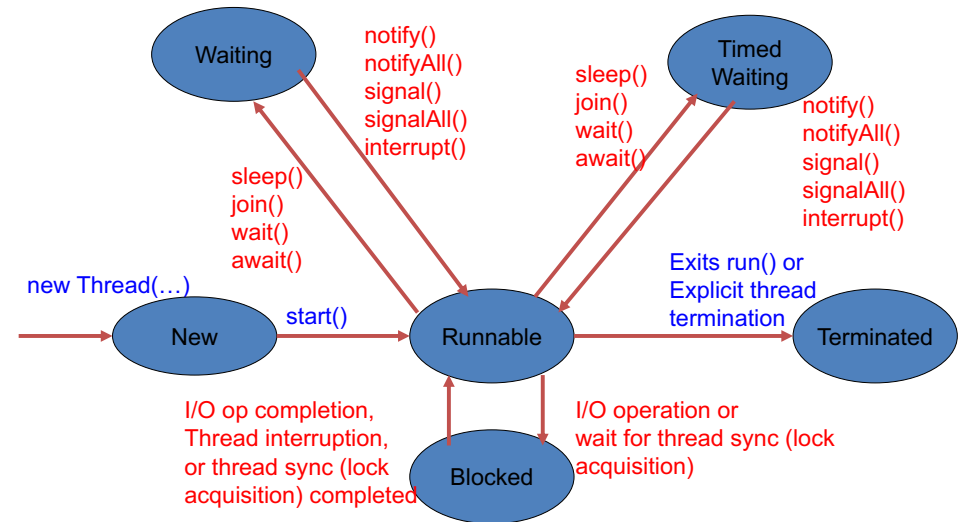
- Used to synchronize (or serialize, or mutually-exclude) multiple threads that access shared data.
- **java.util.concurrent.locks.Lock** interface
  - **ReentrantLock** class: the most commonly-used class for thread synchronization
    - Defines methods that
      - allow threads to access shared data in a synchronized (or serialized, or mutually-excluded) manner.
      - allow you to write atomic code.
- Atomic code is surrounded by **lock()** and **unlock()** method calls.
  - ```
ReentrantLock aLock = new ReentrantLock();  
aLock.lock();  
atomic code (access to shared data)  
aLock.unlock();
```

24

States of a Thread

- Once a thread calls `lock()`,
 - it acquires and owns a **lock** until it calls `unlock()`.
 - No other threads can acquire the lock until it is released by `unlock()`.
 - No other threads can run atomic code until the lock is released with `unlock()`.
- If a thread calls `lock()` when another thread already owns the lock,
 - it goes to the **blocked** state and *gets blocked* (cannot do anything further) until the lock is released.

25



26

How Can a Blocked Thread Run Again?

- JVM's thread scheduler
 - Periodically reactivates all blocked threads so that they can try to acquire the target lock.
 - If the lock is still unavailable, they get blocked again.
 - Detects a release of the target lock (i.e. completion of atomic code).
 - May notify all blocked threads so that one of them can acquire the target lock.
 - May choose one of the blocked threads to acquire the lock.
- Each blocked thread can eventually acquire the target lock when it is available.

27

Coding Idiom for Locking

- Call `unlock()` in a **finally clause**.


```

ReentrantLock aLock = new ReentrantLock();
aLock.lock();
try {
    atomic code (access to a shared variable)
}
finally {
    aLock.unlock();
}
      
```
- `unlock()` is never invoked
 - if `run()` returns in atomic code
 - if atomic code throws an exception
 - A deadlock occurs.
 - Atomic code is locked forever, and no other threads can acquire the lock to run the atomic code.

28

```

• aLock.lock();
  try{
    atomic code
  }
  finally{
    aLock.unlock();
  }

```

DO THIS!


```

• try{
  aLock.lock();
  atomic code
}
finally{
  aLock.unlock();
}

```


DON'T DO THIS!

- Make sure to call lock() **BEFORE** a “try” block.
- If a thread throws an exception in lock(), it will not acquire the lock. However, it will call unlock().
 - lock() can throw an InterruptedException when another thread call interrupt().




Thread synchronization

- Prevents the “withdraw” thread from withdrawing money from “balance” when the “deposit” thread is depositing money to “balance.”
- Prevents the “deposit” thread from depositing money to “balance” when the “withdraw” thread is withdrawing money from “balance.”



ThreadSafeBankAccount.java



```

• private ReentrantLock lock = new ReentrantLock();

• public void deposit(double amount){
  lock.lock();
  try{
    balance += amount; // atomic code
  }finally{
    lock.unlock();
  }
}

• public void withdraw(double amount){
  lock.lock();
  try{
    balance -= amount; // atomic code
  }finally{
    lock.unlock();
  }
}

```

30

Summary: How to Avoid Race Conditions?

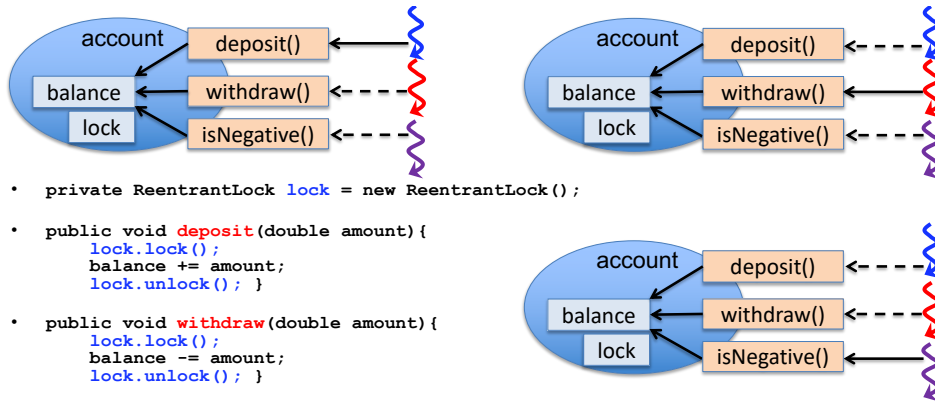
- When multiple threads share and access a variable concurrently,
 - Make sure to **guard the shared variable with a lock**.
 - Identify **ALL read and write logic** to be performed on the variable
 - Surround each of them with lock() and unlock().

```

- e.g., public void deposit(double amount){
  lock.lock();
  try{
    balance += amount; // atomic code (read and write)
  }finally{
    lock.unlock();
  }
}

- public void withdraw(double amount){
  lock.lock();
  try{
    balance -= amount; // atomic code (read and write)
  }finally{
    lock.unlock();
  }
}

```

```
private ReentrantLock lock = new ReentrantLock();
```

```
public void deposit(double amount) {
    lock.lock();
    balance += amount;
    lock.unlock();
}
```

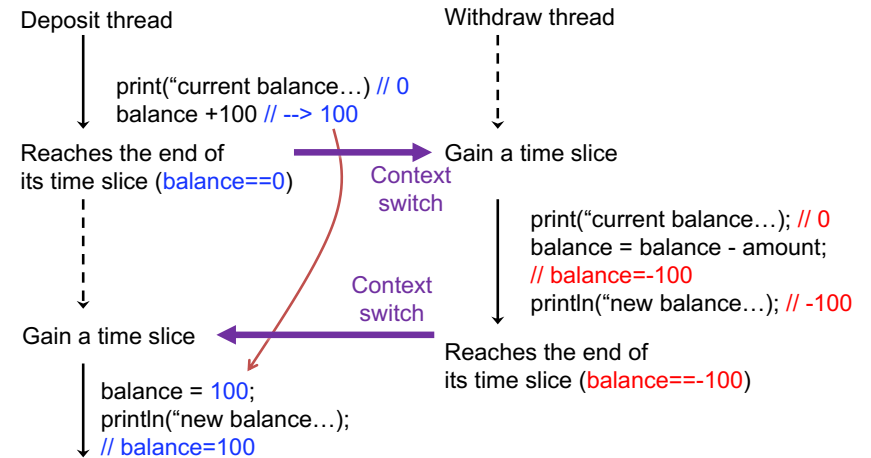
```
public void withdraw(double amount) {
    lock.lock();
    balance -= amount;
    lock.unlock();
}
```

```
public boolean isNegative() {
    lock.lock();
    if (balance < 0) { return true; }
    else { return false; }
    lock.unlock();
}
```

Thread synchronization

- Mutually excludes those 3 threads, so only one of them can run atomic code associated with “lock” at a time.
- It is important to use the **same lock** in **ALL read and write logic** to be performed on “balance.” Otherwise, threads are NOT mutually excluded.

When Could a Context Switch Occur?



Atomicity of Operations for Primitive Types

A context switch can occur

- Across atomic operations.
 - e.g., `int i = 1; i = 2;`
- In a compound operation.
 - e.g., `balance = balance + amount;`

The **read** and **write** operations for primitive data types, except double and long (64-bit) types, are **atomic**.

- An “atomic” operation is transformed to a **single bytecode instruction** for a JVM.
- **No context switches** occur during the execution of a single byte code instruction.

```
int x;
```

- Thread A does: `x=1`; Thread B does: `x=2`;
- An assignment of an int value (write operation) is atomic.
- x contains 1 or 2 depending on which thread performs assignment earlier.
 - x never contains other values (e.g., 0 and 3) or corrupted data.
 - An example of corrupted data
 - » Some part of x (e.g., the first 16-bit of x) comes from Thread A and the remaining part (e.g., the other 16bit of x) comes from Thread B.

Compound Operations

- A *compound* of atomic operations is NOT atomic.
 - `int i; boolean done;`
 - `done = true;` // 2 steps
 - `i = 1;` // 2 steps
 - `if(done)` // 2 steps
 - `i = j;` // 2 steps
 - `j = i + 1;` // 5 steps
 - Reading the value of `i`, reading/loading the value of 1, doing `i+1`, storing the result of `i+1` to a certain memory space, and assigning the result to `j`.
 - `i = i + 1;` // 5 steps
 - `i++` // 5 steps
- A race condition can occur due to a context switch *in between* atomic operations/steps.

Another Example

- ```
public void deposit(double amount){
 balance = balance + amount;
}
```

    - A compound of 5 atomic operations.
    - There are 4 places where race conditions can occur.
  - Thread synchronization enables **serialized** (or **atomic** or **exclusive**) access to a compound operation.
    - Allows **only one thread** to perform a compound op at a time.
- ```
- ReentrantLock aLock = new ReentrantLock();  
aLock.lock();  
try{  
    balance = balance + amount;  
}  
finally{  
    aLock.unlock();  
}
```

An Example Race Condition

- `i = i + 1`
 - A compound of 5 atomic operations.
 - There are 4 places where race conditions can occur.
- Thread synchronization enables **serialized** (or **atomic** or **exclusive**) access to a compound operation.
 - Allows **only one thread** to perform a compound operation at a time.

```
- ReentrantLock aLock = new ReentrantLock();  
aLock.lock();  
try{  
    i = i + 1; // treated as an atomic operation  
}  
finally{  
    aLock.unlock();  
}
```

What about 64-bit Types?

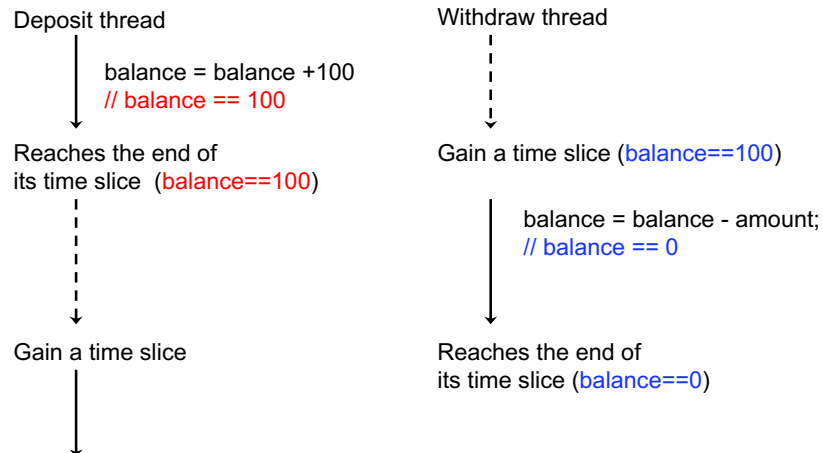
- The **read** and **write** operations for double and long variables are **NOT atomic**.
 - `long x;`
 - Thread A does: `x = 1L;`
 - Thread B does: `x = 2L;`
 - No guarantee that `x` contains 1L or 2L.
 - `x` can contain another value (e.g., 0L or 3L) or corrupted data.
 - `aLongVar = 100L;` // 2+ bytecode instructions
 - `if(aLongVar)` // 2+ bytecode instructions
 - `aLongVar ++` // 5+ bytecode instructions

Atomicity of Operations for Reference Types

- The **read** and **write** operations reference types are **atomic**.
 - A compound of atomic operations
 - e.g., `Foo foo = temp; // 2 byte code instructions`
 - A race condition can occur due to a context switch in between atomic steps.

Consider this Lucky Case

- ```
public void deposit(double amount){
 balance = balance + amount; }
```



# What's Tricky in Thread Programming

- Your test code may or may not be able to detect race conditions.
  - It may not be able to detect race conditions even if you run it a lot of (e.g. a few hundred) times.

## Nested Locking

- ```
class BankAccount {  
    private double balance;  
    private ReentrantLock lock;  
  
    public void deposit(double amount) {  
        lock.lock();  
        balance += amount; // 5 atomic steps  
        if(balance < MIN_BALANCE) // 4 atomic steps  
            subtractPenaltyFee();  
        lock.unlock(); }  
  
    private void subtractPenaltyFee() {  
        balance -= PENALTY; // 5 atomic steps  
        // NO NEED TO SURROUND THIS LINE BY LOCK() and UNLOCK()  
        // because it is called from atomic code.  
    }  
}
```

Thread Reentrancy

```
• class BankAccount {  
    private double balance;  
    private ReentrantLock lock;  
  
    public void deposit(double amount) {  
        lock.lock();  
        balance += amount; // 5 atomic steps  
        if(balance < MIN_BALANCE) // 4 atomic steps  
            subtractPenaltyFee();  
        lock.unlock(); }  
  
    private void subtractPenaltyFee() {  
        lock.lock();  
        balance -= PENALTY; // 5 atomic steps  
        lock.unlock(); } }
```

- This code does not have a deadlock problem.
- A thread can **re-enter** the same lock as far as it already owns the lock.

45

```
• class A{  
    private B b;  
    private ReentrantLock lock;  
    public void a1(){  
        b = new B();  
        lock.lock();  
        b.b1(this); //nested locking  
        lock.unlock();  
    }  
    public void a2(){  
        lock.lock();  
        do something.  
        lock.unlock();  
    }  
}
```

```
• Class B{  
    public void b1(A a){  
        a.a2();  
    }  
}
```

If a thread performs:
A a = new A();
a.a1();
it **re-enters (or re-acquires)** the same lock that it already owns.

- This code does not have a deadlock problem.
- A thread can **re-enter** the same lock as far as it already owns the lock.