


# Recap: Singleton Design Pattern

- Guarantee that a class has only one instance.
  - c.f. CS680 lecture note
- ```
public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;

    // Factory method to return the singleton instance
    public static Singleton getInstance(){
        if(instance==null)
            instance = new Singleton();
        return instance;
    }
}
```
- This code is NOT thread-safe; race conditions can occur.




- When multiple threads call `getInstance()` concurrently, they share instance.
- ```
public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;

    // Factory method to return the singleton instance
    public static Singleton getInstance(){
        if(instance==null)           // Read   3 steps
            instance = new Singleton(); // Write 2+ steps
        return instance;             // Read   2 steps
    }
}
```

1

2

# Concurrent Singleton Design Pattern



- ```
public class Singleton{
    private Singleton(){};
    private static Singleton instance = null;
    private static Singleton instance = null;
    // Write 2 steps
    // Can assume this to be
    // performed atomically }
}
```
- JVM completes all **initial value assignments on all static data fields** BEFORE using a class or creating class instances.
  - `instance` has been initialized before a thread(s) call `getInstance()`
  - You can assume this write logic is performed atomically.
    - No need to worry about race conditions here.

3

- Guarantee that a class has only one instance.

- ```
public class ConcurrentSingleton{
    private Singleton(){};
    private static Singleton instance = null;
    private static ReentrantLock lock = new ReentrantLock();

    // Factory method to create or return the singleton instance
    public static Singleton getInstance(){
        lock.lock();
        try{
            if(instance==null){ instance = new Singleton(); }
            return instance;
        }finally{
            lock.unlock();
        }
    }
}
```

4

# HW 10

- The Singleton class is not thread-safe.
  - Race conditions can occur if you do not guard the `instance` variable with a lock. Explain a potential race condition in which more than one instances are created.
    - Use a diagram like in a previous slide.
- Submit a thread-safe version of it (`ConcurrentSingleton`)
  - Define a lock in Singleton. Use the lock in `getInstance()`
    - Use try-finally blocks: Always do this in all subsequent HWs.
  - Create multiple extra threads and have them call `getInstance()`
    - Make sure that only one instance is created.
      - Use `System.out.println(Singleton.getInstance())`

## Regular and Static Locks

- ```
public class Foo{
    ReentrantLock lock = new ReentrantLock();
    static ReentrantLock sLock = new ReentrantLock(); }
```
- A regular lock is created and used on an instance-by-instance basis.
  - Different instances of `Foo` have different locks (i.e., different instances of `ReentrantLock`).
- A static lock is created and used on a per-class basis.
  - All instances of `Foo` share a single lock (`sLock`).

- Deadline: Oct 25 (Thu) midnight

## Exercise: Regular and Static Locks

- ```
public class Foo{
    private ReentrantLock lock = new ReentrantLock();
    private static ReentrantLock sLock = new ReentrantLock();

    public void a() {...}
    public void b() {...}
    public void syncA() {lock.lock(); ... lock.unlock();}
    public void syncB() {lock.lock(); ... lock.unlock();}

    public static void sA() {...}
    public static void sB() {...}
    public static void sSyncA() {sLock.lock(); ... sLock.unlock();}
    public static void sSyncB() {sLock.lock(); ... sLock.unlock();} }
```
- `x = new Foo(); y = new Foo();`
- Two threads call...
  - `x.a()` and `x.a()`: no synchronization (no mutual exclusion) for the two threads
  - `x.a()` and `x.b()`: no synchronization
  - `x.a()` and `x.syncA()`: no synchronization
  - `x.syncA()` and `x.syncA()`: Synchronization (mutual exclusion)
  - `y.syncA()` and `y.syncB()`: Synchronization
  - `x.syncA()` and `y.syncA()`: No synchronization
  - `x.syncA()` and `y.syncB()`: No synchronization

- ```

public class Foo{
    private ReentrantLock    lock = new ReentrantLock();
    private static ReentrantLock sLock = new ReentrantLock();

    public void a(){...}
    public void b(){...}
    public void syncA(){lock.lock(); ... lock.unlock();}
    public void syncB(){lock.lock(); ... lock.unlock();}

    public static void sA(){...}
    public static void sB(){...}
    public static void sSyncA(){sLock.lock(); ... sLock.unlock();}
    public static void sSyncB(){sLock.lock(); ... sLock.unlock();} }

```
- ```

x = new Foo(); y = new Foo();

```
- Two threads call...
  - x.a() and Foo.sA(): No synchronization for the two threads
  - x.syncA() and Foo.sA(): No synchronization
  - Foo.sA() and Foo.sA(): No synchronization
  - Foo.sA() and Foo.sB(): No synchronization
  - x.syncA() and Foo.sSyncA(): No synchronization
  - Foo.sSyncA() and Foo.sSyncA(): Synchronization
  - Foo.sSyncA() and Foo.sSyncB(): Synchronization
  - x.sSyncA() and y.sSyncB(): Synchronization
    - This is not grammatically wrong, but write `Foo.sSyncA()` instead of `x.sSyncA()`

## Thread.sleep()

- ```

Thread t = new Thread( new FooRunnable() );
t.start();
try{
    t.sleep(1000);
}catch(InterruptedException e){...}

```
- It looks like an extra thread (t) will sleep.
- However, the main thread will actually sleep
  - because `sleep()` is a **static method** of Thread.
    - `Thread.sleep()`: Causes the currently executing thread to sleep (temporarily cease execution) for the specified number of milliseconds
- DO NOT write `t.sleep(...)`. It's misleading and error-prone.
- ALWAYS WRITE `Thread.sleep(...)`.
  - Make sure to do this in HW 9.

## RunnableInterruptiblePrimeGenerator

```

class InterruptiblePrimeGenerator extends PrimeGenerator {
    public void generatePrimes(){
        for (long n = from; n <= to; n++){
            if( Thread.interrupted() ){
                System.out.println("Stopped");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); } } }

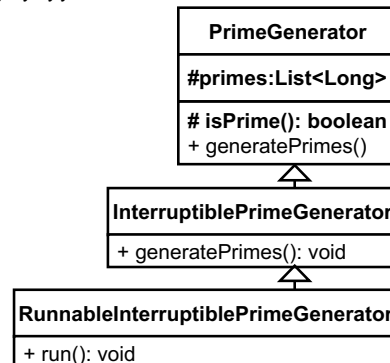
```

```

class RunnableInterruptiblePrimeGenerator
    extends InterruptiblePrimeGenerator
    implements Runnable {

    public void run(){
        generatePrimes(); } }

```



## interrupt(), isInterrupted() and interrupted()

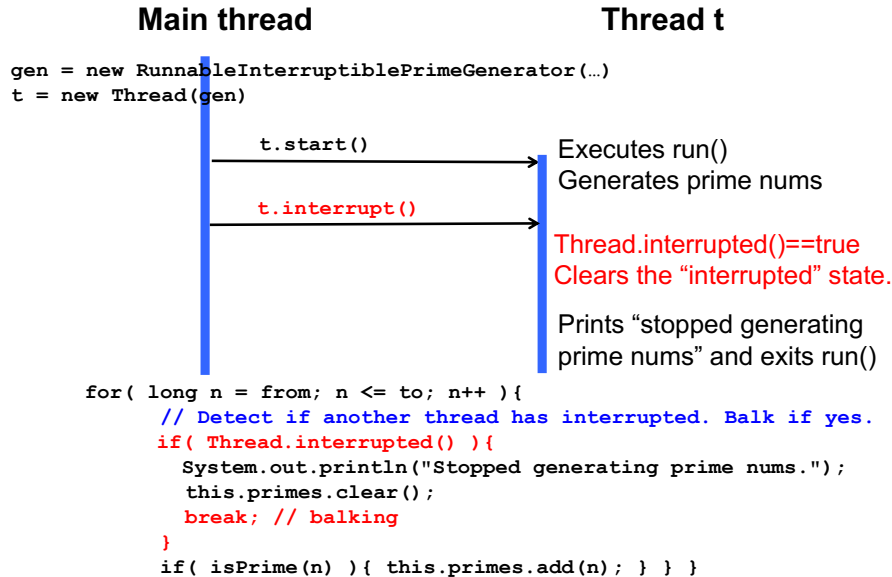
- ```

public class Thread{
    public void interrupt();
    public boolean isInterrupted();
    public static boolean interrupted();
}

```
- Each thread (Thread instance) has the “interrupted” (boolean) state.
- interrupt()**
  - Interrupts **this** thread and changes its “interrupted” state.
    - `aThread = new Thread(...); aThread.start();`  
`aThread.interrupt();`
- isInterrupted()**
  - Returns true if **this** thread has been interrupted.
    - `aThread = new Thread(...); aThread.start();`  
`if( aThread.isInterrupted() ){...}`
  - Does not change the “interrupted” state of the thread.
- interrupted()**
  - Returns true if the currently-executed thread has been interrupted.
  - Clears the “interrupted” state (true → false) if true is returned.

## Thread Interruption != Thread Termination

- `interrupt()` **NEVER** terminate a thread.
  - It simply change the “interrupted” state
    - to help/trigger a thread termination.

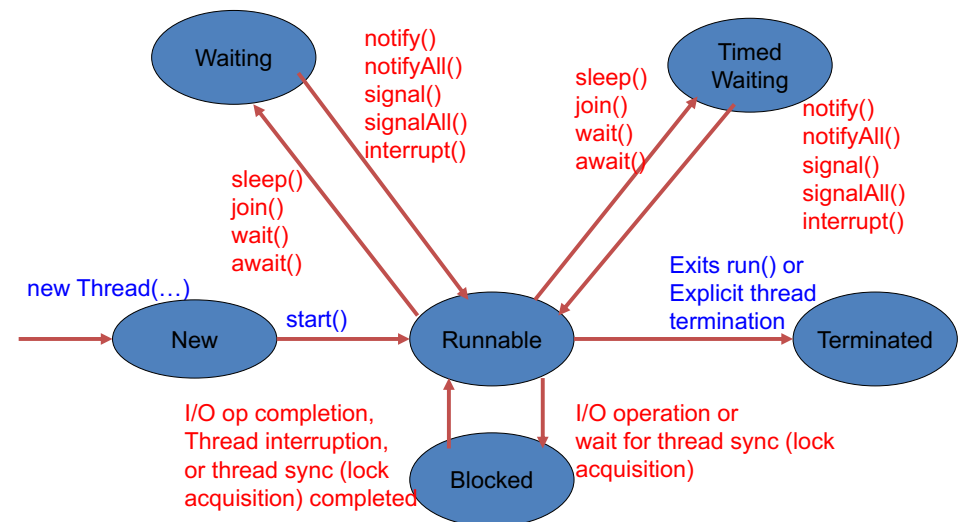


13

## What Happens When `interrupt()` is Called on a Thread?

- If the soon-to-be-terminated thread is in the Runnable state, `interrupt()` changes its “interrupted” state to be true.
- If the soon-to-be-terminated thread is in the *Waiting* or *Blocked* state, it throws an `InterruptedException`.

## States of a Thread



16

## RunnableInterruptiblePrimeGenerator

- In fact, RunnableInterruptiblePrimeGenerator is NOT thread-safe. Race conditions can occur.

```
class InterruptiblePrimeGenerator extends PrimeGenerator {
    public void generatePrimes(){
        for (long n = from; n <= to; n++){
            if( Thread.interrupted() ){ // 2 steps
                System.out.println("Stopped");
                this.primes.clear();
                break;
            }
            if( isPrime(n) ){ this.primes.add(n); } }
    }

    class RunnableInterruptiblePrimeGenerator
        extends InterruptiblePrimeGenerator
        implements Runnable {

        public void run(){
            generatePrimes(); } }
}
```

## Thread.interrupt()

- ```
public void interrupt(){
    ...
    synchronized(...){ // Acquire a lock in Thread
        ...
        interrupt0();    // native method (atomic)
    }
}
```
- ```
public static boolean interrupted(){
    return currentThread().isInterrupted(true); // native method
                                                // (atomic)
}
```
- interrupt() and interrupted() are thread-safe.
  - isInterrupted() is thread-safe as well.
  - c.f. Java source code (e.g. greppcode.com)
- However, *client code* of interrupted() is NOT guaranteed to be thread-safe.

## Solution: Locking and Balking

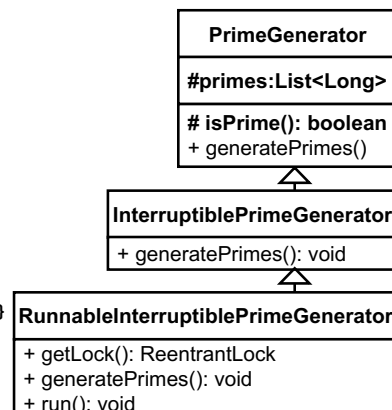
```
class RunnableInterruptiblePrimeGenerator
    extends InterruptiblePrimeGenerator
    implements Runnable {

    private final ReentrantLock lock = new ReentrantLock();

    public ReentrantLock getLock(){
        return lock; }

    public void generatePrimes(){
        for (long n = from; n <= to; n++){
            lock.lock();
            if( Thread.interrupted() ){
                System.out.println("Stopped");
                this.primes.clear();
                break;
            }
            lock.unlock();
            if(isPrime(n)) {this.primes.add(n); } }

    public void run(){
        generatePrimes(); } }
}
```



- Main thread (client of RunnableInterruptiblePrimeGenerator)
  - ```
RunnableInterruptiblePrimeGenerator gen =
    new RunnableInterruptiblePrimeGenerator();
Thread aThread = new Thread(gen); aThread.start();

gen.getLock().lock();
aThread.interrupt();
gen.getLock().unlock();
```
- This code uses two locks.
  - One in Thread
  - One in RunnableInterruptiblePrimeGenerator

## HW 11

- Revise `RunnableInterruptiblePrimeGenerator.java` to be thread-safe.
  - c.f. HW 10, in which you work on a thread-safe version of `RunnableCancelablePrimeGenerator.java`
- Deadline: Oct 25 (Thu) midnight

## Hybridization of the Two Approaches?

- Can we implement a **responsive** thread termination that uses **only 1 lock**?

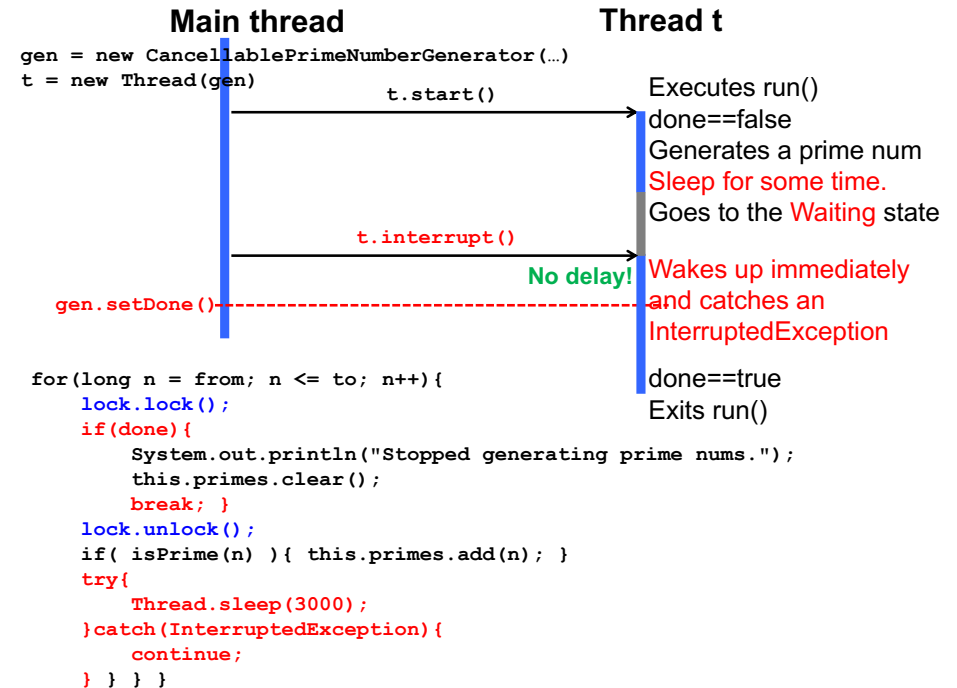
## Explicit Thread Termination

- Flag-based
  - Pros:
    - Uses **1 lock (faster)**
  - Cons:
    - Program responsiveness may be lower.
      - if a flag-flipping (e.g. `done==false → true`) happens when a thread-to-be-terminated is in the Waiting or Blocked state.
- Interruption-based
  - Pros
    - Higher program responsiveness
      - `interrupt()` can immediately wake up a thread-to-be-terminated that is in the Waiting or Blocked state
  - Cons
    - Uses **2 locks (slower)**

## 2-Step Thread Termination ("Graceful" Thread Termination)

## 2-Step Thread Termination

- Primarily takes the flag-based approach.
  - A thread-to-be-terminated periodically checks a flag.
- Let the “terminator” thread call `interrupt()` before flipping the flag’s state (i.e., before calling `setDone()`)

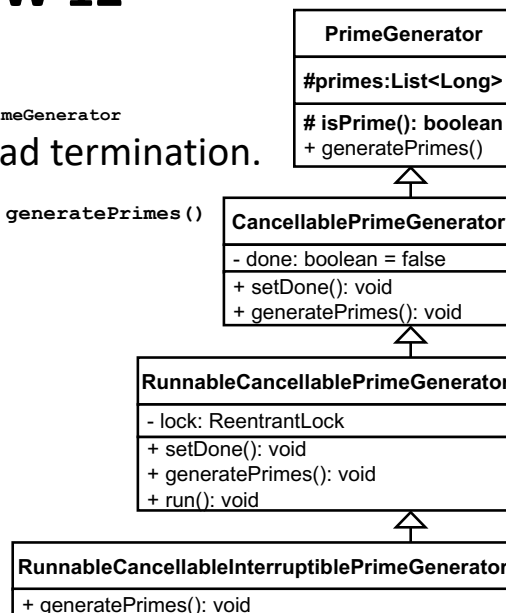


## HW 12

- Define `RunnableCancellableInterruptiblePrimeGenerator` to perform 2-step thread termination.

– Re-define (or override) `generatePrimes()`

- Deadline: Oct 29 (Tue) midnight




## 2-Step Thread Termination is Effective if...

- A thread-to-be-terminated may be in the **Waiting or Blocked state** when a “terminator” thread tries to terminate it.
  - Performing an I/O operation.
    - e.g., reading/writing data from/to a file, waiting for an incoming data on a socket, sending data to a remote app.
  - Waiting for a lock acquisition
    - Has called `lock()` on a lock, but the lock is not available yet.
  - Has called `sleep()`, `join()`, etc.

## What Happens When `interrupt()` is Called on a Thread?

- If a soon-to-be-terminated thread is in the Runnable state, `interrupt()` changes its “interrupted” state to be true.
- If the soon-to-be-terminated thread is in the *Waiting* or *Blocked* state, it raises an `InterruptedException`.

## States of a Thread



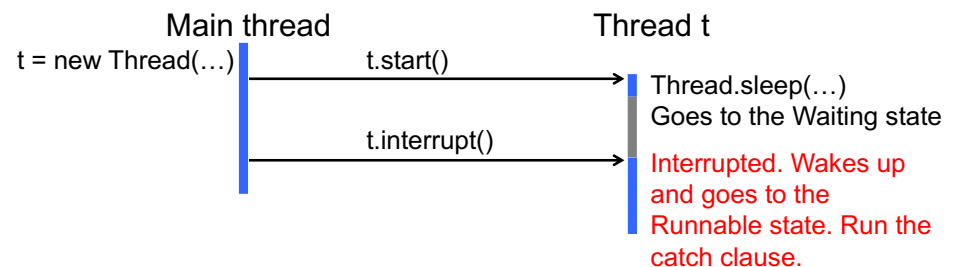
30

## InterruptedException

- Some methods in Java API throws `InterruptedException`.
  - They can respond to a thread interruption by throwing an `InterruptedException`.
  - `Thread.sleep()`
  - `Thread.join()`
  - `ReentrantLock.lockInterruptibly()`
  - `BlockingQueue.put()/take()`
  - `Condition.await()`
  - I/O operations
  - These methods can be long-running and **interruptible**.

## Thread.sleep()

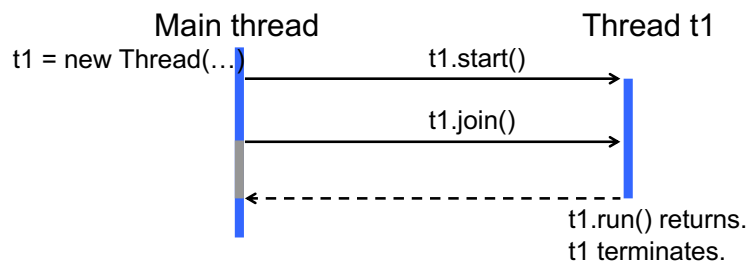
- `sleep()` lets the *currently-executed thread* to sleep for a specified time period.
- `interrupt()` interrupts a sleeping thread.
  - Wakes up the thread and force `sleep()` to throw an `InterruptedException`.
- ```
try{
    Thread.sleep(60000);
}catch(InterruptedException e){
    // Write thread termination (shutdown) logic here.
}
```






## Thread.join()

- `join()` lets the *currently-executed thread* to wait/sleep until another thread terminates (i.e., until another thread returns `run()`).
- `interrupt()` can interrupt a waiting/sleeping thread.
  - Force `join()` to throw an `InterruptedException`.



## Thread.join()


- `join()` lets the *currently-executed thread* to wait/sleep until another thread terminates (i.e., until another thread returns `run()`).
- `interrupt()` can interrupt a waiting/sleeping thread.
  - Force `join()` to throw an `InterruptedException`.



## Condition.await()

- `await()` lets the currently-executed thread wait/sleep until another thread wakes it up with `signal()/signalAll()`.
- `interrupt()` can interrupt a waiting/sleeping thread.
  - Allows `await()` to acquire a lock and forces it to throw an `InterruptedException`

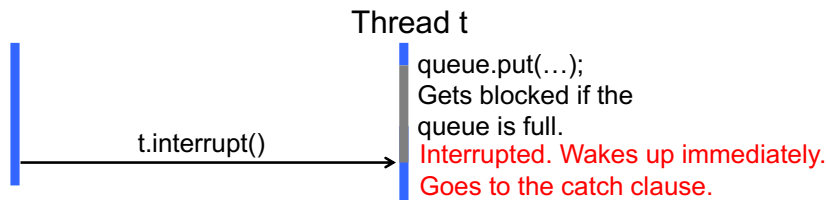
```
withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        try{
            // waiting for the balance to exceed 0
            sufficientFundsCondition.await();
        }catch(InterruptedException e){
            //Do something
        }
    }
    belowUpperLimitFundsCondition.signalAll();
    balance -= amount;
    lock.unlock();
}
```



```
withdraw(double amount){
    lock.lock();
    while(balance <= 0){
        try{
            // waiting for the balance to exceed 0
            sufficientFundsCondition.await();
        }catch(InterruptedException e){
            //Do something; e.g., balk with a
            // "break" statement.
        }
    }
    ...
}
```

# BlockingQueue

- `interface BlockingQueue<E> extends Queue<E>`
  - Adds A Queue that additionally supports operations that
    - wait for the queue to become non-empty when retrieving an element
    - wait for space to become available in the queue when storing an element.
- Several impls: `ArrayBlockingQueue`, `LinkedBlockingQueue`, etc.
  - `put()` and `take()` are *blocking* methods.
    - `put()`: Add an element to a queue as the last element.
    - `take()`: Get the first element in the queue.
  - They can respond to a thread interruption by throwing an `InterruptedException`.



# Where did the synchronized Keyword go?

- Java still has the `synchronized` keyword.
  - ```
public synchronized void foo(){  
    // The entire method body is atomic.  
}
```
  - ```
public void foo(){  
    // non-atomic code here  
    synchronized(this){  
        // atomic code here  
    }  
    // non-atomic code here }
```
- Implicit locking
  - No need to create a `ReentrantLock` and call `lock()` and `unlock()`.
- When a thread enters a synchronized method/block, it tries to acquire the (implicit) lock that `this` instance maintains.
  - Instance-by-instance locking
- Code gets tricky/dirty to use multiple locks in a single class.

# Thread Termination

- Thread creation is a no brainer.
- Thread termination requires your careful attention.
  - No methods available in `Thread` to directly terminate threads like `terminate()`.
    - Do: 2-step termination
  - Why not?
    - Different programmers/apps need different termination policies.
      - Notify on-going thread termination to other threads?
      - Raise exception(s) in addition to `InterruptedException`?
      - What to do for the data maintained by a thread being terminated?
    - Java allows you to flexibly craft your own termination policy.

## Explicit locking

```
ReentrantLock aLock = new ReentrantLock()  
public void foo(){  
    aLock.lock();  
    // atomic code  
    aLock.unlock(); }
```

- Arbitrary locking scope.
- Clean code even if a class uses multiple locks.
- Extra functionalities
  - e.g., `getQueueLength()`: returns the # of waiting threads.
  - `tryLock()`: acquires a lock only if it is not held by another thread.
- The catch is... it's **VERY easy to forget calling `unlock()`**.
  - Must call `unlock()` in a finally clause.

- Implicit locking with the “synchronized” keyword
  - A thread can call notify() and notifyAll() even if it has not acquired a lock.
    - An IllegalMonitorStateException is thrown.
- Explicit locking
  - This error/bug never occurs.
    - ```
ReentrantLock lock = new ReentrantLock();  
Condition cond = lock.newCondition();  
lock.lock();  
...  
cond.signalAll();
```