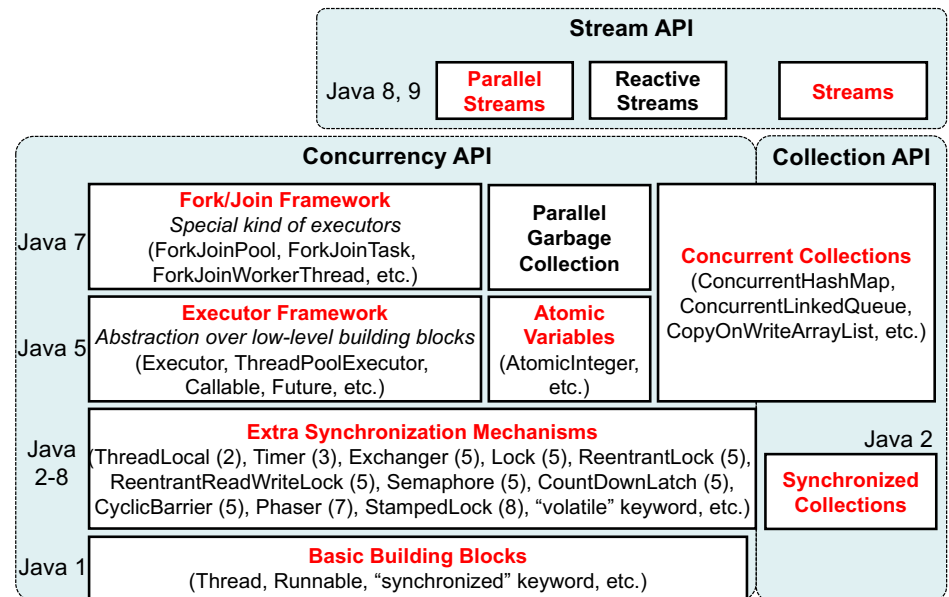


Executor Framework

Executor Framework

- An **abstraction layer** atop low-level concurrency primitives
 - Focuses on **task execution** on threads
 - Decouples task execution (on threads) from task submission (to threads) to **make task execution configurable**.
 - Introduced in Java 5 (2004)
 - Enhanced further in subsequent versions
 - Implemented in `java.util.concurrent`.

Concurrency API in Java




Tasks, Threads and Executor

- **Tasks**
 - Logical units of work
 - e.g., prime number generation, access counting for files, banking (deposit/withdrawal/wire transfer of money), file caching, file crawling, file indexing, etc.
- **Threads**
 - Mechanism to run tasks *concurrently*.
- **Executor**
 - Is the primary abstraction for task execution
 - **Thread** is NOT anymore.

Executor

- `public interface Executor{
 void execute(Runnable task); }`
- `Runnable`'s `run()` implements a task.
- **Producers**: submit tasks
- **Consumers**: execute tasks
- Makes task execution configurable.



6


Task Execution Policies

- The Executor framework allows you to specify and customize the *execution policy* for tasks.
 - “What, where, when and how” of task execution.
 - In which thread will tasks be executed?
 - In what order should tasks be executed (FIFO, LIFO, priority-based ordering)?
 - How many tasks may run concurrently.
 - How many tasks may be queued pending execution?
 - If a task has to be rejected because an application is overloaded, which task should be selected as the victim? How should the application be notified?
 - What actions should be taken before or after executing a task?

7


Thread Pool

- A key component for task execution.
- A set of **pre-created “worker” threads** that will be used for future task execution
- Each worker thread
 - Gets and executes a task if it is available in the queue.
 - Goes to the Waiting state, if no tasks are available in the queue, until a producer submits the next task.



8

- Benefits of using a thread pool
 - Can eliminate runtime overhead to create threads
 - Can bound the maximum number of threads (i.e., the max amount of resource utilization)
 - Running too many threads (i.e., consuming too much resources) will result in a crash of operating system.



9

Executors

- A utility class for `Executor` objects

- Defines **static factory methods** to create an executor with a particular thread pool.

```

- static ExecutorService newFixedThreadPool(int n)
    • Fixed-size thread pool

- ExecutorService executor = Executors.newFixedThreadPool(2);
  executor.execute( new PrimeNumberGenerator(1L, 500000L) );
  executor.execute( new PrimeNumberGenerator(500001L, 1000000L) );

- Thread t1, t2;
  t1 = new Thread( new PrimeNumberGenerator(1L, 500000L) );
  t2 = new Thread( new PrimeNumberGenerator(500001L, 1000000L) );
  t1.start();
  t2.start();

```

10

- ```


- static ExecutorService newCachedThreadPool()
 • Variable-size (not fixed-size) thread pool.
 – Uses previously created “idle” threads if they are available.
 – Creates a new thread if no idle threads are available.
 – Idle threads are terminated and removed from the pool after they are not used for 60 seconds.

 • Pros:
 – Can minimize the number of tasks in the queue.
 – Can minimize the number of threads (resource consumption)

 • Cons: No cap for the number of threads in the pool.

 • Useful to handle a number of short-lived (lightweight) tasks

```



12

- Static factory methods

- ```

- static ExecutorService newFixedThreadPool(int n)
    • Fixed-size thread pool.

- static ScheduledExecutorService newScheduledThreadPool(int n)
    • Fixed-size thread pool that supports delayed and periodic task execution.

- static ExecutorService newSingleThreadExecutor()
    • A pool that operates only one thread.

- static ScheduledExecutorService newSingleThreadScheduledExecutor()
    • A single-threaded pool that supports delayed and periodic task execution.

```

11

- ```


- static ExecutorService newWorkStealingPool(int parallelism)
 • Variable-size thread pool with a cap for the # of threads.
 – Parallelism specifies the cap for the # of threads.

 • Each worker thread
 – Has its own “primary” queue and gets the next task from the queue.
 – “Steals” a task from another queue if no tasks are available in its primary queue.
 – Dies after being idle for some time.

 • Pros:
 – Each queue requires less thread synchronization.
 – Can minimize the # of tasks in a queue and bound the # of worker threads.

 • Cons: No guarantee about the order task execution


```



13

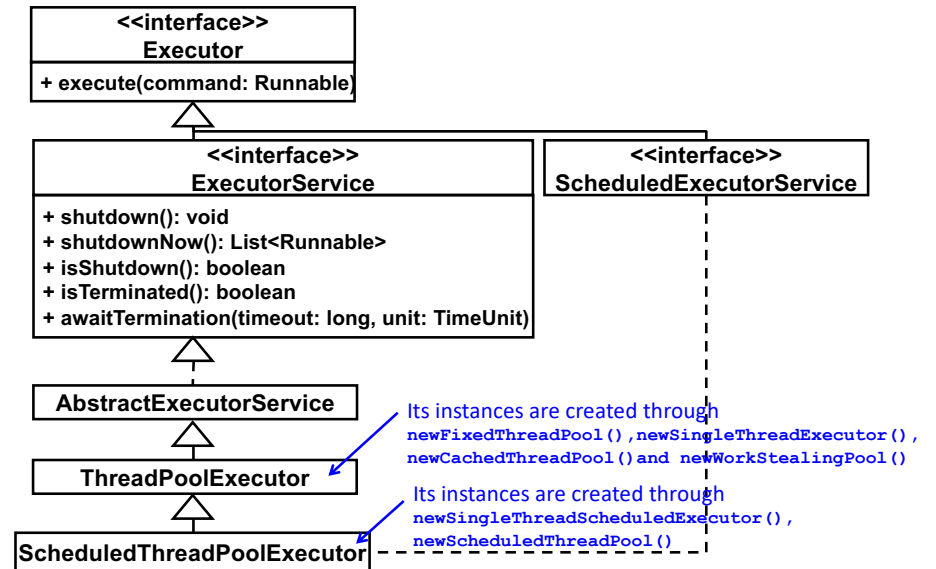
- `static ExecutorService newWorkStealingPool()`

– Obtains the number of available CPU cores by calling `availableProcessors()` and invokes the previous version of `newWorkStealingPool()`



14

## ExecutorService



15

## Termination of Executor

- Methods to terminate an executor

### – `shutdown()`


- Rejects new tasks to get in
  - Throws a `RejectedExecutionException`
- Allows previously submitted tasks to complete
  - Tasks being executed and tasks in the queue

### – `shutdownNow()`

- Rejects new tasks to get in
- Removes all tasks from the queue and returns them
- Tries to stop the tasks that are being executed.
  - Call `interrupt()` on each worker thread

- A task can be stopped if it checks `Thread.interrupted()` or catches `InterruptedException` to exit `run()`.

– Otherwise, it may not be stopped



- 3 states of an executor


– **Running**

– **Shutting down**

- Once `shutdown()` or `shutdownNow()` is called.
- `isShutdown()` returns true.

– **Terminated**

- Once all tasks have been completed or stopped.
- `isTerminated()` returns true.



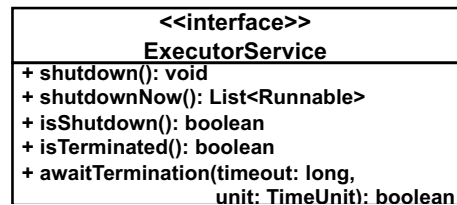
- Use `awaitTermination()` if you wait for an executor to be terminated.
  - It blocks until the executor is terminated or the timeout occurs.
  - It returns true if the executor is terminated or false otherwise.

```

- executor.shutdown();
 executor.awaitTermination(Long.MAX_VALUE, TimeUnit.SECONDS);
 doSomething();

- executor.shutdown();
 if(!executor.awaitTermination(60, TimeUnit.SECONDS)){
 shutdownNow();
 if(!executor.awaitTermination(60, TimeUnit.SECONDS)){
 doErrorHandling();
 }
 }

```



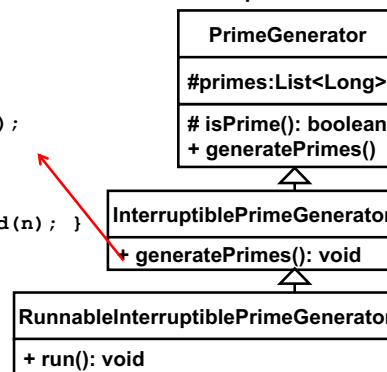
## RunnableInterruptiblePrimeGenerator

- Detect an interruption from another thread to stop generating prime numbers.

```

- for (long n = from; n <= to; n++){
 if(Thread.interrupted()){
 System.out.println("Stopped");
 this.primes.clear();
 break;
 }
 if(isPrime(n)){ this.primes.add(n); }
}

```



## Sample Code: RunnableInterruptiblePrimeGenExecutorTest.java

```

• RunnableInterruptiblePrimeGenerator r1, r2;
 r1 = new RunnableInterruptiblePrimeGenerator(1L, 500000L);
 r2 = new RunnableInterruptiblePrimeGenerator(500001L, 1000000L);

 ExecutorService executor = Executors.newFixedThreadPool(2);

 executor.execute(r1);
 executor.execute(r2);

 executor.shutdown();

 //executor.shutdownNow();
 // Calls interrupt() on each prime gen thread. An
 // interruption is caught by Thread.interrupted() in
 // RunnableInterruptiblePrimeGenerator's run().

 executor.awaitTermination(...);

 r1.getPrimes().forEach(...);
 r2.getPrimes().forEach(...);

```

19

- If you use `shutdown()`,
  - Two runnable tasks generate all primes.
    - gen1 generated 41538 prime numbers.
    - gen2 generated 36960 prime numbers.
- If you use `shutdownNow()`,
  - Two runnable tasks cancel prime generation
    - Stopped generating prime numbers due to a thread interruption.
    - Stopped generating prime numbers due to a thread interruption.
    - gen1 generated 0 prime numbers.
    - gen2 generated 0 prime numbers.